Logging is not (distributed) Tracing

Questions (to the audience) free!

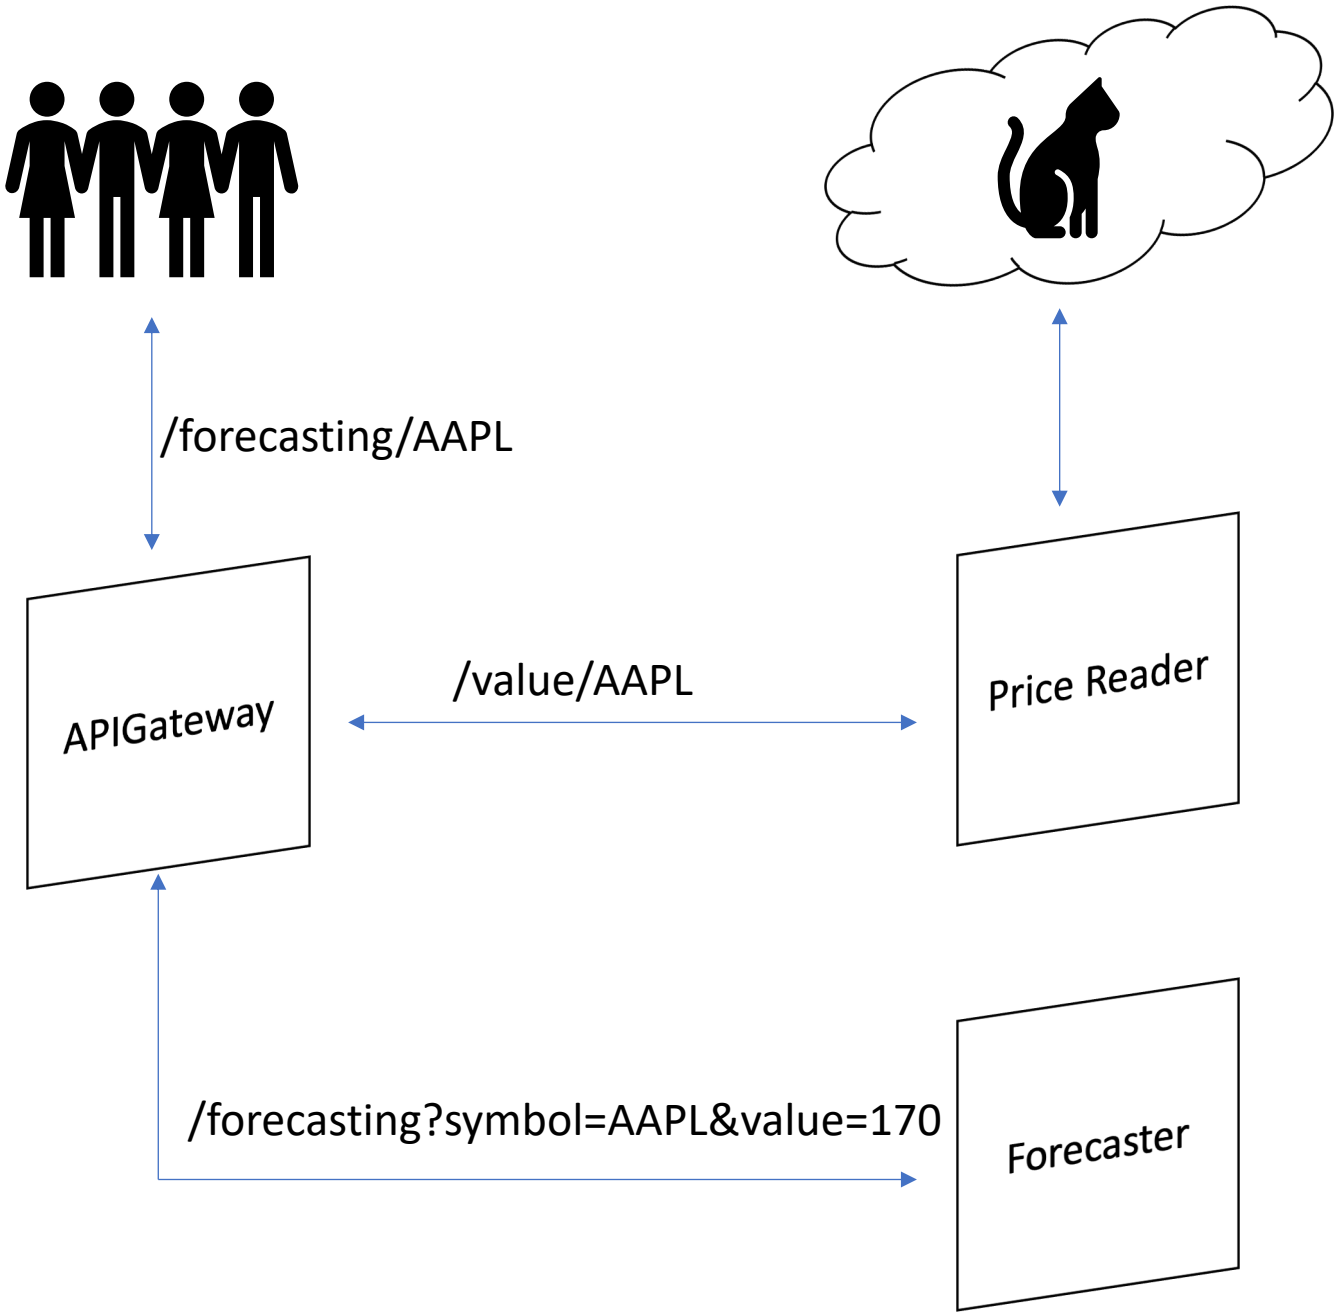
/forecasting/AAPL
APIGateway
/value/AAPL
Price Reader
/forecasting?symbol=AAPL&value=170
Forecaster

Logging is

Storing execution information for later analysis

Logging...

printf(“Error\n”);

std::cout << “Error\n”;

What?

Define your goal and audience

Decide what to log

Silence  Noise

What?

Context

Severity levels and Categories

Not only for troubleshooting

Not only for humans

Where

Log to a file (and rotate logs) and, maybe, centralized (Graylog, ELK Stack…)

How?

Use a logging library

spdlog::error(“Error");

spdlog

https://github.com/gabime/spdlog/wiki/1.-QuickStart

Distributed Tracing is

a method used to profile and monitor applications, especially those built using a microservices architecture

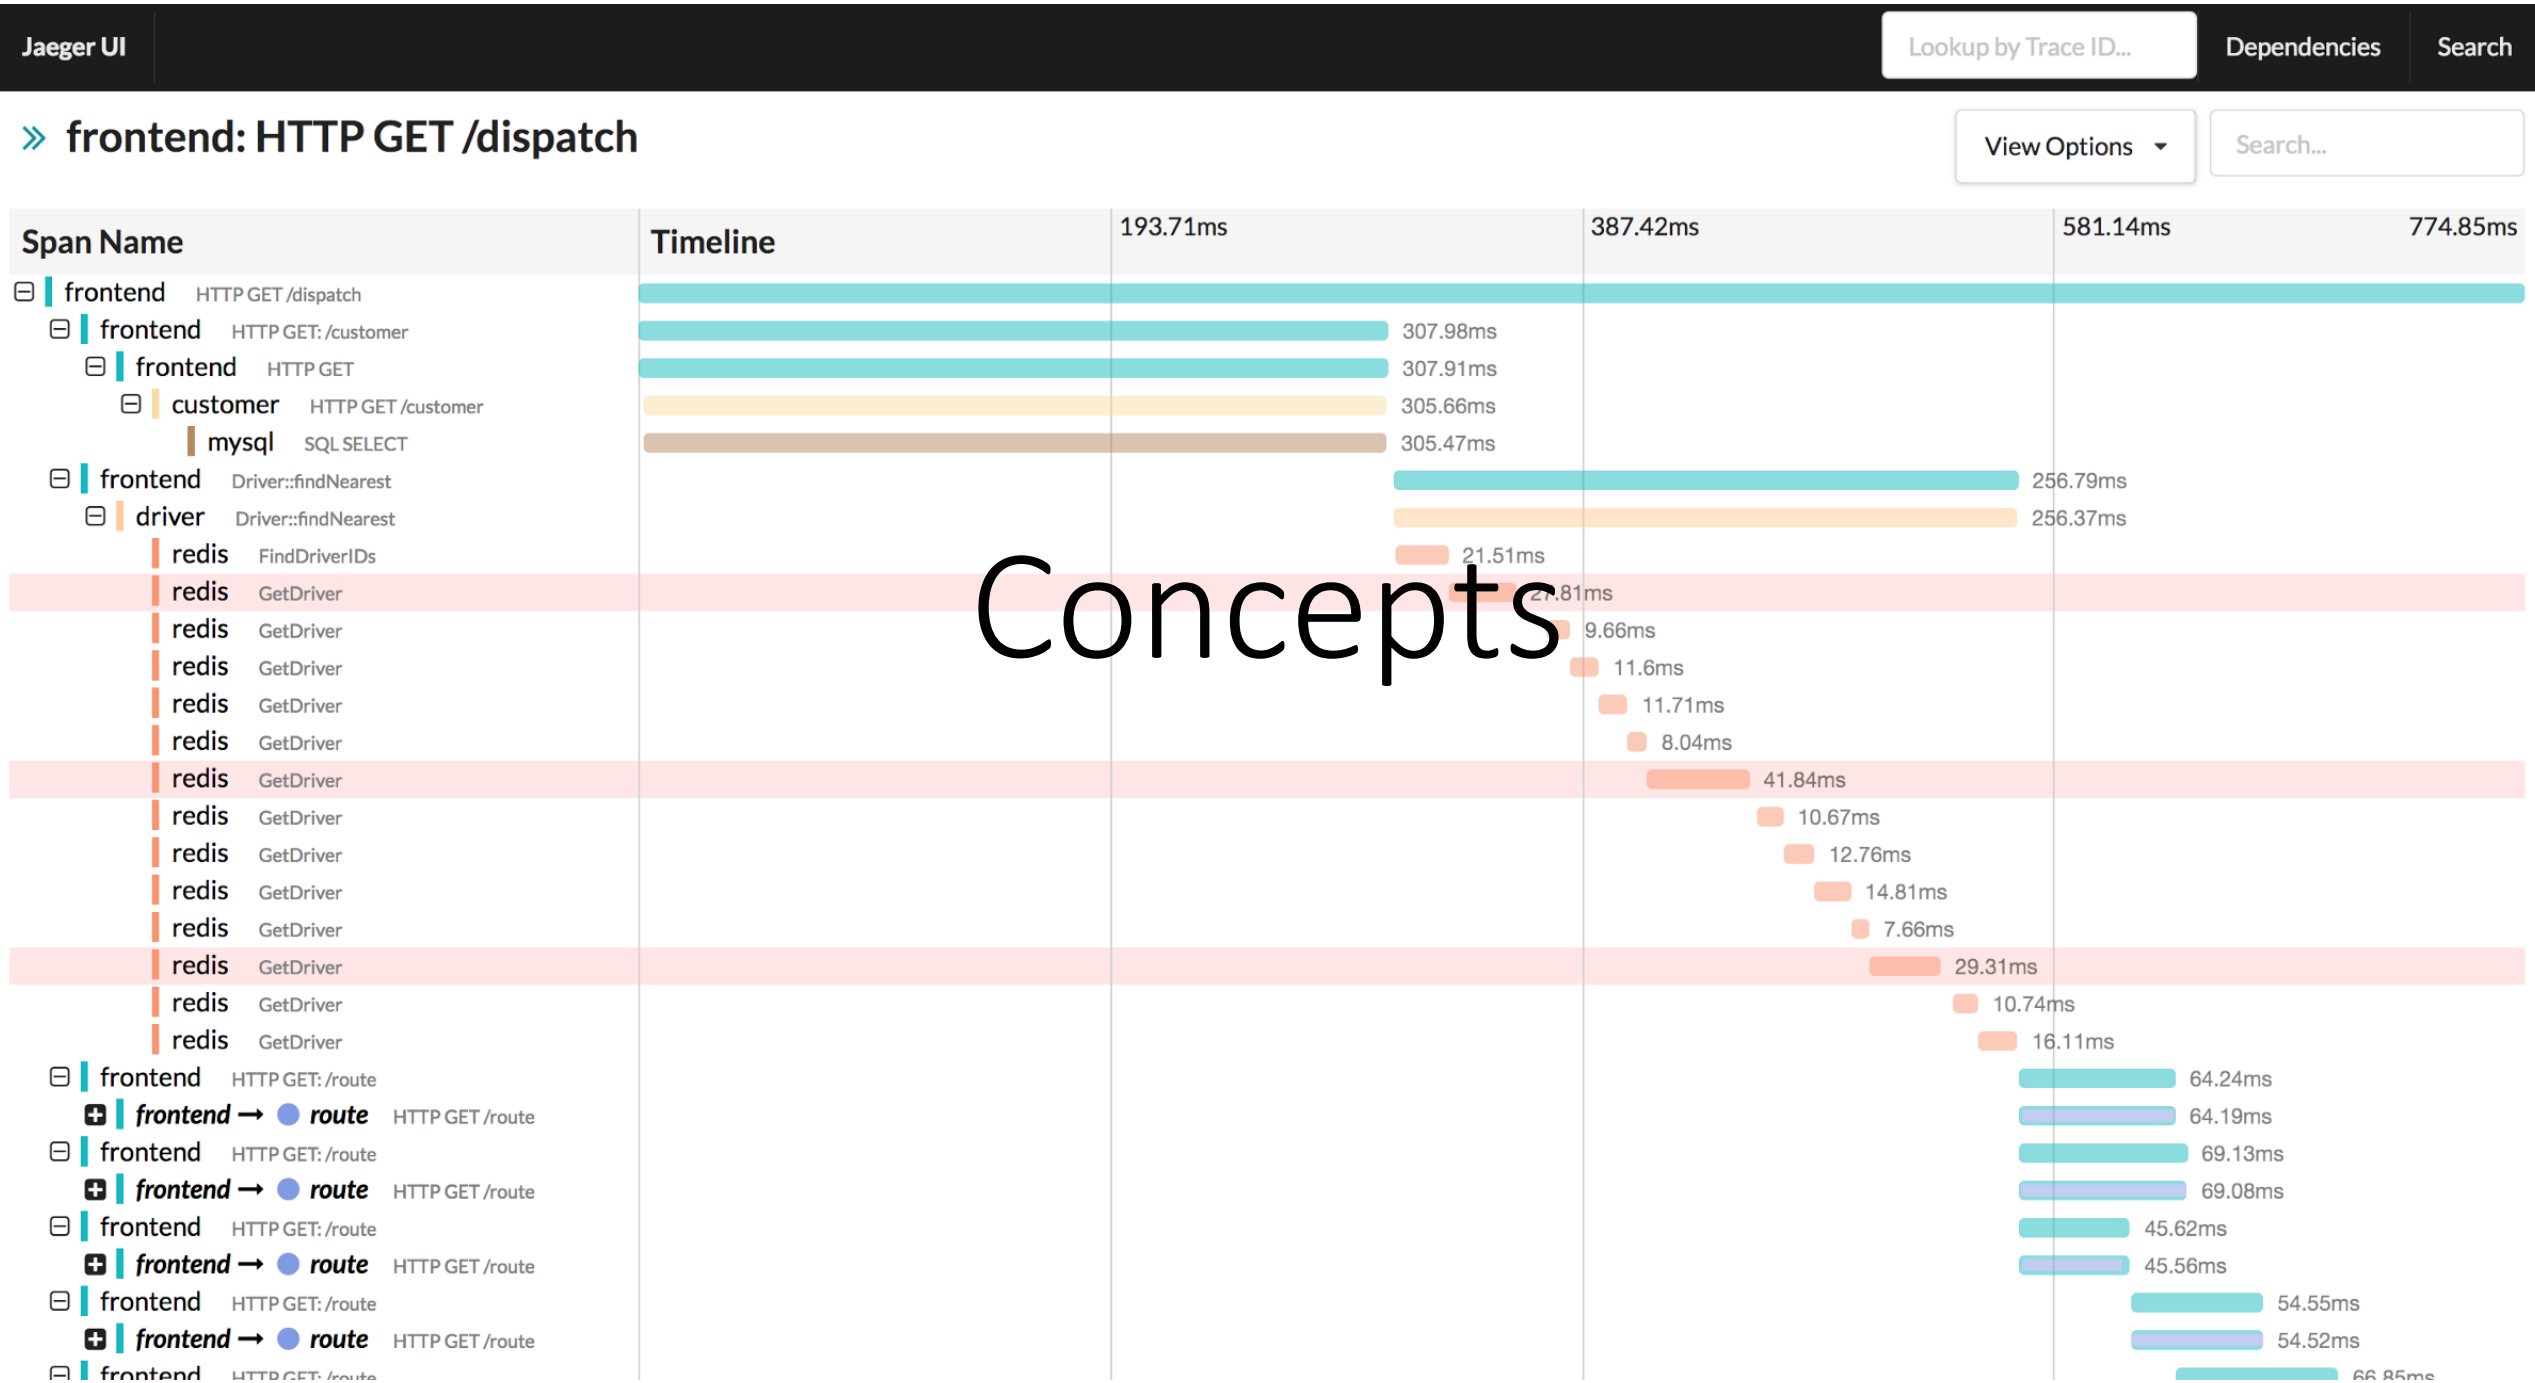
Jaeger UI
Lookup by Trace ID...
Dependencies
Search
frontend: HTTP GET /dispatch
View Options
Search...
Span Name
Timeline
193.71ms
387.42ms
581.14ms
774.85ms
frontend HTTP GET /dispatch
frontend HTTP GET: /customer
307.98ms
frontend HTTP GET
307.91ms
customer HTTP GET /customer
305.66ms
mysql SQL SELECT
305.47ms
frontend Driver::findNearest
256.79ms
driver Driver::findNearest
256.37ms
redis FindDriverIDs
21.51ms
redis GetDriver
redis GetDriver
9.66ms
redis GetDriver
11.6ms
redis GetDriver
11.71ms
redis GetDriver
8.04ms
redis GetDriver
41.84ms
redis GetDriver
10.67ms
redis GetDriver
12.76ms
redis GetDriver
14.81ms
redis GetDriver
7.66ms
redis GetDriver
29.31ms
redis GetDriver
10.74ms
redis GetDriver
16.11ms
frontend HTTP GET: /route
64.24ms
frontend → route HTTP GET /route
64.19ms
frontend HTTP GET: /route
69.13ms
frontend → route HTTP GET /route
69.08ms
frontend HTTP GET: /route
45.62ms
frontend → route HTTP GET /route
45.56ms
frontend HTTP GET: /route
54.55ms
frontend → route HTTP GET /route
54.52ms
Concepts

Tracing

Tracing tracks the progression of a single user request as it is handled by other services that make up an application.

Trace

A directed acyclic graph (DAG) of Spans

Span

A span represents a single operation in a trace

a Span has

Name, start and finish timestamps

Tags: [{ key, value }]

Logs: [{ key, value, timestamp }]

and a SpanContext

Information needed to refer to a distinct Span across a process boundary and zero or more Baggage Items as [{ key, value }] that cross process boundaries

and References

ChildOf: a Span is the child of a parent Span

FollowsFrom: the new Span is independent of the parent Span

Tigres, Leones

OpenTracing (CNCF) vs OpenCensus (Google)

Jaeger

An OpenTracing Tracer (and an OpenCensus Exporter)

Why?

C++ client

Demo

https://github.com/david-antiteum/logging

The End